

Basic data manipulation

Lecture 3

Louis SIRUGUE




CPES 2 - Fall 2022



Quick reminder

1. Distributions

- The **distribution** of a variable documents all its possible values and how frequent they are






- We can describe a distribution with:



Quick reminder

1. Distributions

- The **distribution** of a variable documents all its possible values and how frequent they are






- We can describe a distribution with:
 - Its **central tendency**



Quick reminder

1. Distributions

- The **distribution** of a variable documents all its possible values and how frequent they are



- We can describe a distribution with:
 - Its **central tendency**
 - And its **spread**



Quick reminder

2. Central tendency

- The **mean** is the sum of all values divided by the number of observations

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

- The **median** is the value that divides the (sorted) distribution into two groups of equal size

$$\text{Med}(x) = \begin{cases} x_{[\frac{N+1}{2}]} & \text{if } N \text{ is odd} \\ \frac{x_{[\frac{N}{2}]} + x_{[\frac{N}{2}+1]}}{2} & \text{if } N \text{ is even} \end{cases}$$

3. Spread

- The **standard deviation** is square root of the average squared deviation from the mean

$$\text{SD}(x) = \sqrt{\text{Var}(x)} = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2}$$

- The **interquartile range** is the difference between the maximum and the minimum value from the middle half of the distribution

$$\text{IQR} = Q_3 - Q_1$$


Quick reminder

4. Inference

- In Statistics, we view variables as a given realization of a **data generating process**
 - Hence, the **mean** is what we call an **empirical moment**, which is an **estimation...**
 - ... of the **expected value**, the **theoretical moment** of the DGP we're interested in
- To know how confident we can be in this estimation, we need to compute a **confidence interval**

$$\left[\bar{x} - t_{n-1, 97.5\%} \times \frac{SD(x)}{\sqrt{n}}; \bar{x} + t_{n-1, 97.5\%} \times \frac{SD(x)}{\sqrt{n}} \right]$$

- It gets **larger** as the **variance** of the distribution of x increases
- And gets **smaller** as the **sample size** n increases



Warm up practice

05:00

- 1) Import the `ligue1.csv` dataset and store it in an object called `fb`
- 2) Create a subset of this dataset containing only matches that took place at 13h
- 3) Print the number of matches in this subset and compute the average attendance
- 4) Redo the same exercise on matches that took place at 20h45

You've got 5 minutes!

Solution

1) Import the `ligue1.csv` dataset and store it in an object called `fb`

```
fb <- read.csv("C:/User/Documents/ligue1.csv", encoding = "UTF-8")
```

2) Create a subset of this dataset containing only matches that took place at 13h

```
sub13 <- fb[fb$Time == "13:00", ]
```

3) Print the number of matches in this subset and compute the average attendance

```
nrow(sub13)
```

```
## [1] 32
```

```
mean(sub13$Attendance)
```

```
## [1] NA
```


Solution

- When there are **missing values** in a vector, the **mean** function returns **NA**
 - We need to set the **na.rm** option to **TRUE**

3) Print the number of matches in this subset and compute the average attendance

```
mean(sub13$Attendance, na.rm = T)
```

```
## [1] 19038
```

4) Redo the same exercise on matches that took place at 20h45

```
sub2045 <- fb[fb$Time == "20:45", ]  
nrow(sub2045)
```

```
## [1] 29
```

```
mean(sub2045$Attendance, na.rm = T)
```

```
## [1] 36418.64
```



Today we learn how to manipulate data

1. The dplyr package

- 1.1. Packages
- 1.2. Basic functions
- 1.3. `group_by()` and `summarise()`

2. Merge and reshape

- 2.1. Merge and append data
- 2.2. Reshape data

3. A few words on learning R

- 3.1. When it doesn't work the way you want
- 3.2. Where to find help
- 3.3. When it doesn't work at all

4. Wrap up!



Today we learn how to manipulate data

1. The dplyr package

1.1. Packages

1.2. Basic functions

1.3. `group_by()` and `summarise()`



1. The `dplyr` package

1.1. Packages

- So far we only used functions that are directly available in R
 - But there are tons of **user-created functions** out there that can make your life so much easier
 - These functions are shared in what we call **packages**
- Packages are **bundles of functions** that R users put at the disposal of other R users
 - Packages are **centralized** on the [Comprehensive R Archive Network \(CRAN\)](#)
 - To **download** and install a CRAN package you can simply use **`install.packages()`**
- All the functions of the `dplyr` grammar are gathered in the **`dplyr` package**
 - We can download these functions and make them ready to use with the `install.packages()` function

```
install.packages("dplyr") # Requires an internet connection
```

- The tidyverse package is **now installed** on your computer
 - You won't have to do it again



1. The `dplyr` package

1.1. Packages

- The `dplyr` package is now **on your computer**, but it is **not loaded in R**

```
ls("package:dplyr")
```

```
## Error in as.environment(pos): no item called "package:dplyr" on the search list
```

- You need to use the **`library()`** command to load it

```
library(dplyr)  
ls("package:dplyr")[1:5]
```

```
## [1] "%>%"      "across"    "add_count" "add_count_" "add_row"
```

- But even though the package is permanently installed, it is **loaded only for your current session**
 - Each time you start a **new R session**, you'll have to load the packages you need with **`library()`**

1. The `dplyr` package

1.2. Basic functions

`dplyr` is a **grammar** of data manipulation providing very **user-friendly functions** to handle the most common **data manipulation** tasks:

- `mutate()`: add/modify variables
- `select()`: keep/drop variables (columns)
- `filter()`: keep/drop observations (rows)
- `arrange()`: sort rows according to the values of given variable(s)
- `summarise()`: aggregate the data into descriptive statistics



- A very handy **operator** to use with the **dplyr** grammar is the **pipe** `%>%`
 - You can basically read `a %>% b()` as "*apply function b() to object a*"
 - With this operator you can easily **chain the operations** you apply to an object



1. The `dplyr` package

1.2. Basic functions

```
fb
```

```
#  
#  
#  
#  
#  
#
```

##	Wk	Day	Date	Time	Home	xG	Score	xG.1	Away	Attendance	...
## 1	1	Fri	2021-08-06	21:00	Monaco	2.0	1-1	0.3	Nantes	7500	...
## 2	1	Sat	2021-08-07	17:00	Lyon	1.4	1-1	0.8	Brest	29018	...
## 3	1	Sat	2021-08-07	21:00	Troyes	0.8	1-2	1.2	Paris S-G	15248	...
## 4	1	Sun	2021-08-08	13:00	Rennes	0.6	1-1	2.0	Lens	22567	...
## 5	1	Sun	2021-08-08	15:00	Bordeaux	0.7	0-2	3.3	Clermont Foot	18748	...
## 6	1	Sun	2021-08-08	15:00	Strasbourg	0.4	0-2	0.9	Angers	23250	...
## 7	1	Sun	2021-08-08	15:00	Nice	0.8	0-0	0.2	Reims	18030	...
## 8	1	Sun	2021-08-08	15:00	Saint-Étienne	2.1	1-1	1.3	Lorient	20461	...
## 9	1	Sun	2021-08-08	17:00	Metz	0.7	3-3	1.4	Lille	15551	...
...



1. The `dplyr` package

1.2. Basic functions

```
fb %>%  
  select(Home, xG, Score, xG.1, Away)           # Keep/drop certain columns  
#  
#  
#  
#  
#  
#
```

	Home	xG	Score	xG.1	Away
## 1	Monaco	2.0	1-1	0.3	Nantes
## 2	Lyon	1.4	1-1	0.8	Brest
## 3	Troyes	0.8	1-2	1.2	Paris S-G
## 4	Rennes	0.6	1-1	2.0	Lens
## 5	Bordeaux	0.7	0-2	3.3	Clermont Foot
## 6	Strasbourg	0.4	0-2	0.9	Angers
## 7	Nice	0.8	0-0	0.2	Reims
## 8	Saint-Étienne	2.1	1-1	1.3	Lorient
## 9	Metz	0.7	3-3	1.4	Lille
...



1. The `dplyr` package

1.2. Basic functions

```
fb %>%  
  select(Home, xG, Score, xG.1, Away) %>%  
  mutate(home_winner = xG > xG.1)  
  
# Keep/drop certain columns  
# Create a new variable  
#  
#  
#  
#
```

##	Home	xG	Score	xG.1	Away	home_winner
## 1	Monaco	2.0	1-1	0.3	Nantes	TRUE
## 2	Lyon	1.4	1-1	0.8	Brest	TRUE
## 3	Troyes	0.8	1-2	1.2	Paris S-G	FALSE
## 4	Rennes	0.6	1-1	2.0	Lens	FALSE
## 5	Bordeaux	0.7	0-2	3.3	Clermont Foot	FALSE
## 6	Strasbourg	0.4	0-2	0.9	Angers	FALSE
## 7	Nice	0.8	0-0	0.2	Reims	TRUE
## 8	Saint-Étienne	2.1	1-1	1.3	Lorient	TRUE
## 9	Metz	0.7	3-3	1.4	Lille	FALSE
...



1. The `dplyr` package

1.2. Basic functions

```
fb %>%  
  select(Home, xG, Score, xG.1, Away) %>%  
  mutate(home_winner = xG > xG.1) %>%  
  filter(Home == "Rennes")  
  
# Keep/drop certain columns  
# Create a new variable  
# Keep/drop certain rows  
#  
#  
#
```

##	Home	xG	Score	xG.1	Away	home_winner
## 1	Rennes	0.6	1-1	2.0	Lens	FALSE
## 2	Rennes	0.9	1-0	0.5	Nantes	TRUE
## 3	Rennes	1.0	0-2	0.5	Reims	TRUE
## 4	Rennes	2.4	6-0	0.3	Clermont Foot	TRUE
## 5	Rennes	0.8	2-0	1.4	Paris S-G	FALSE
## 6	Rennes	1.5	1-0	0.6	Strasbourg	TRUE
## 7	Rennes	3.8	4-1	1.1	Lyon	TRUE
## 8	Rennes	3.1	2-0	0.7	Montpellier	TRUE
## 9	Rennes	0.8	1-2	0.6	Lille	TRUE



1. The `dplyr` package

1.2. Basic functions

```
fb %>%  
  select(Home, xG, Score, xG.1, Away) %>%  
  mutate(home_winner = xG > xG.1) %>%  
  filter(Home == "Rennes") %>%  
  arrange(-xG)  
  
# Keep/drop certain columns  
# Create a new variable  
# Keep/drop certain rows  
# Sort rows  
#  
#
```

##		Home	xG	Score	xG.1	Away	home_winner
## 1		Rennes	3.8	4-1	1.1	Lyon	TRUE
## 2		Rennes	3.3	6-0	0.4	Bordeaux	TRUE
## 3		Rennes	3.3	6-1	0.9	Metz	TRUE
## 4		Rennes	3.1	2-0	0.7	Montpellier	TRUE
## 5		Rennes	2.7	2-0	0.3	Brest	TRUE
## 6		Rennes	2.6	4-1	0.4	Troyes	TRUE
## 7		Rennes	2.4	6-0	0.3	Clermont Foot	TRUE
## 8		Rennes	1.9	2-3	2.9	Monaco	FALSE
## 9		Rennes	1.7	2-0	0.3	Angers	TRUE
	



1. The `dplyr` package

1.2. Basic functions

```
fb %>%  
  select(Home, xG, Score, xG.1, Away) %>%  
  mutate(home_winner = xG > xG.1) %>%  
  filter(Home == "Rennes") %>%  
  arrange(-xG) %>%  
  summarise(expected_wins = mean(home_winner),  
            expected_goals = sum(xG))
```

Keep/drop certain columns
Create a new variable
Keep/drop certain rows
Sort rows
Aggregate into statistics
#

```
##      expected_wins expected_goals  
## 1          0.8421053           36.6
```



1. The `dplyr` package

1.2. Basic functions

- Here are two very **handy functions** to use within `mutate()`

`ifelse`

```
fb %>%  
  select(Home, Attendance) %>%  
  mutate(att_bin = ifelse(Attendance > 10000,  
                           "Large",  
                           "Low")  
         ) %>% head()
```

##	Home	Attendance	att_bin
## 1	Monaco	7500	Low
## 2	Lyon	29018	Large
## 3	Troyes	15248	Large
## 4	Rennes	22567	Large
## 5	Bordeaux	18748	Large
## 6	Strasbourg	23250	Large

`case_when`

```
fb %>%  
  select(Home, xG, xG.1, Away) %>%  
  mutate(xWin = case_when(xG > xG.1 ~ "Home",  
                           xG == xG.1 ~ "Draw",  
                           xG < xG.1 ~ "Away")  
         ) %>% head()
```

##	Home	xG	xG.1	Away	xWin
## 1	Monaco	2.0	0.3	Nantes	Home
## 2	Lyon	1.4	0.8	Brest	Home
## 3	Troyes	0.8	1.2	Paris S-G	Away
## 4	Rennes	0.6	2.0	Lens	Away
## 5	Bordeaux	0.7	3.3	Clermont Foot	Away
## 6	Strasbourg	0.4	0.9	Angers	Away



1. The `dplyr` package

1.3. `group_by()` and `summarise()`

- With `group_by()` you can perform **computations separately** for the different **categories of a variable**

```
fb %>%  
  select(Wk, Home, xG) %>%  
  mutate(all.xG = mean(xG)) %>%  
  head(10)
```

##	Wk	Home	xG	all.xG
## 1	1	Monaco	2.0	1.473421
## 2	1	Lyon	1.4	1.473421
## 3	1	Troyes	0.8	1.473421
## 4	1	Rennes	0.6	1.473421
## 5	1	Bordeaux	0.7	1.473421
## 6	1	Strasbourg	0.4	1.473421
## 7	1	Nice	0.8	1.473421
## 8	1	Saint-Étienne	2.1	1.473421
## 9	1	Metz	0.7	1.473421
## 10	1	Montpellier	0.5	1.473421

```
fb %>%  
  select(Wk, Home, xG) %>%  
  group_by(Home) %>%  
  mutate(home.xG = mean(xG)) %>%  
  head(6)
```

##	#	A tibble:	6 x 4		
##	#	Groups:	Home [6]		
##		Wk	Home	xG	home.xG
##		<int>	<chr>	<dbl>	<dbl>
##	1	1	Monaco	2	1.69
##	2	1	Lyon	1.4	2.07
##	3	1	Troyes	0.8	1.21
##	4	1	Rennes	0.6	1.93
##	5	1	Bordeaux	0.7	1.23
##	6	1	Strasbourg	0.4	1.73



1. The `dplyr` package

1.3. `group_by()` and `summarise()`

- It is particularly **useful with `summarise()`**
 - `summarise` keeps the grouping variable
 - and computes **statistics for each category**

```
fb %>%
  group_by(Wk) %>%
  summarise(n = n(),
            tot_xG = sum(xG)+sum(xG.1),
            avg_WG = tot_xG/n) %>%
  head(4)
```

```
## # A tibble: 4 x 4
##       Wk       n tot_xG avg_WG
##   <int> <int>   <dbl>   <dbl>
## 1     1     10    23.4     2.34
## 2     2     10    26.6     2.66
## 3     3     10    25.7     2.57
## 4     4     10    30.4     3.04
```

`mutate()` \neq `summarise()`

- **`mutate()`** takes an operation that converts:
 - **A vector into another vector**
- **`summarise()`** takes an operation that converts:
 - **A vector into a value**

Ungrouping

- **`group_by()`** applies to all subsequent operations
- To cancel its effect you must **`ungroup()`** the data

```
fb %>%
  group_by(Wk) %>%
  mutate(test = mean(xG)) %>%
  ungroup() %>%
  ...
```

Practice

10:00

- 1) Start from the `fb` dataset and keep only the variables `Home`, `Score` and `Away`
- 2) Use the `separate()` function from `tidyr` to split the `Score` variable into `home_score` and `away_score`

```
data.frame(x = "a_b") %>%  
  separate(x, c("x", "y"), "_")
```

```
##      x y  
## 1 a b
```

- 3) Convert these two variables into numeric vectors
- 4) Create a variable named `winner` that takes the values `Home`, `Draw` and `Away` depending on the score
- 5) Use `group_by()` and `summarise()` to compute the percentage of draws, home wins and away wins

You've got 10 minutes!

Solution

1) Start from the `fb` dataset and keep only the variables `Home`, `Score` and `Away`

```
fb %>%  
  select(Home, Score, Away) %>%  
  head(2)
```

```
##      Home Score  Away  
## 1 Monaco  1-1 Nantes  
## 2   Lyon  1-1  Brest
```

2) Use the `separate()` function from `tidyr` to split the `Score` variable into `home_score` and `away_score`

```
fb %>%  
  select(Home, Score, Away) %>%  
  separate(Score, c("home_score", "away_score"), "-") %>%  
  head(2)
```

```
##      Home home_score away_score  Away  
## 1 Monaco         1         1 Nantes  
## 2   Lyon         1         1  Brest
```

Solution

3) Convert these two variables into numeric vectors

4) Create a variable named **winner** that takes the values **Home**, **Draw** and **Away** depending on the score

```
fb %>%  
  select(Home, Score, Away) %>%  
  separate(Score, c("home_score", "away_score"), "-") %>%  
  mutate(home_score = as.numeric(home_score),  
         away_score = as.numeric(away_score),  
         winner = case_when(home_score < away_score ~ "Away",  
                             home_score == away_score ~ "Draw",  
                             home_score > away_score ~ "Home")) %>%  
  head()
```

##	Home	home_score	away_score	Away	winner
## 1	Monaco	1	1	Nantes	Draw
## 2	Lyon	1	1	Brest	Draw
## 3	Troyes	1	2	Paris S-G	Away
## 4	Rennes	1	1	Lens	Draw
## 5	Bordeaux	0	2	Clermont Foot	Away
## 6	Strasbourg	0	2	Angers	Away

Solution

5) Use `group_by()` and `summarise()` to compute the percentage of draws, home wins and away wins

```
fb %>%
  select(Home, Score, Away) %>%
  separate(Score, c("home_score", "away_score"), "-") %>%
  mutate(home_score = as.numeric(home_score),
         away_score = as.numeric(away_score),
         winner = case_when(home_score < away_score ~ "Away",
                             home_score == away_score ~ "Draw",
                             home_score > away_score ~ "Home")) %>%
  group_by(winner) %>%
  summarise(pct = 100 * (n() / nrow(fb)))
```

```
## # A tibble: 3 x 2
##   winner    pct
##   <chr>   <dbl>
## 1 Away    30.5
## 2 Draw    26.8
## 3 Home    42.6
```



Overview

1. The dplyr package ✓

- 1.1. Packages
- 1.2. Basic functions
- 1.3. group_by() and summarise()

2. Merge and reshape

- 2.1. Merge and append data
- 2.2. Reshape data

3. A few words on learning R

- 3.1. When it doesn't work the way you want
- 3.2. Where to find help
- 3.3. When it doesn't work at all

4. Wrap up!



Overview

1. The dplyr package ✓

- 1.1. Packages
- 1.2. Basic functions
- 1.3. `group_by()` and `summarise()`

2. Merge and reshape

- 2.1. Merge and append data
- 2.2. Reshape data



2. Merge and reshape

2.1. Merge and append data

- Research projects often imply to **combine data** from different sources
 - Either to **add observations** (append rows)
 - Either to **add variables** (merge columns)

Dataset 1 on attainment

country	year	share_tertiary
FRA	2015	44.68760
GBR	2015	49.94341
USA	2015	46.51771


2. Merge and reshape

2.1. Merge and append data

- Research projects often imply to **combine data** from different sources
 - Either to **add observations** (append rows)
 - Either to **add variables** (merge columns)

Dataset 1 on attainment

country	year	share_tertiary
FRA	2015	44.68760
GBR	2015	49.94341
USA	2015	46.51771



Dataset 2 on attainment

country	year	share_tertiary
ITA	2015	25.14996
ESP	2015	40.95978

2. Merge and reshape

2.1. Merge and append data



- Research projects often imply to **combine data** from different sources
 - Either to **add observations** (append rows)
 - Either to **add variables** (merge columns)

Dataset 1 on attainment

country	year	share_tertiary
FRA	2015	44.68760
GBR	2015	49.94341
USA	2015	46.51771

Dataset on spending

country	year	share_gdp
FRA	2015	3.398
USA	2015	3.207
RUS	2015	1.843



Dataset 2 on attainment

country	year	share_tertiary
ITA	2015	25.14996
ESP	2015	40.95978



2. Merge and reshape

2.1. Merge and append data: The `bind_rows()` function

```
read.csv("attainment_FR_UK_US.csv")
```

```
##   country year share_tertiary
## 1     FRA 2015      44.68760
## 2     GBR 2015      49.94341
## 3     USA 2015      46.51771
```

```
read.csv("attainment_IT_SP.csv")
```

```
##   country year share_tertiary
## 1     ITA 2015      25.14996
## 2     ESP 2015      40.95978
```

```
attainment <- read.csv("attainment_FR_UK_US.csv") %>%
  bind_rows(read.csv("attainment_IT_SP.csv"))
attainment
```

```
##   country year share_tertiary
## 1     FRA 2015      44.68760
## 2     GBR 2015      49.94341
## 3     USA 2015      46.51771
## 4     ITA 2015      25.14996
## 5     ESP 2015      40.95978
```

Variables in the two datasets should be the same:

- **Same name**
- **Same class**



2. Merge and reshape

2.1. Merge and append data: `*_join()` functions

- Join functions all work the same way:
 - A **dataset A** with a **variable X** and other variables
 - A **dataset B** with a **variable X** and other variables
 - X is the common variable, so datasets will be **joined by X**

The 4 main join functions

Function	For X in A & B	For X in A only	For X in B only	Summary
A %>% left_join(B, by = "X")	Kept	Kept	Dropped	Only keeps what's in A
A %>% right_join(B, by = "X")	Kept	Dropped	Kept	Only keeps what's in B
A %>% inner_join(B, by = "X")	Kept	Dropped	Dropped	Only keeps what's common
A %>% full_join(B, by = "X")	Kept	Kept	Kept	Keeps everything



2. Merge and reshape

⚠ Beware of NAs! ⚠

- When you have **values** of X that are **not common** to both datasets
 - Any other join than the `inner_join()` will **generate NAs**

```
attainment %>% full_join(read.csv("spending.csv"), by = "country")
```

```
##   country year.x share_tertiary year.y share_gdp
## 1     FRA   2015    44.68760    2015     3.398
## 2     GBR   2015    49.94341      NA        NA
## 3     USA   2015    46.51771    2015     3.207
## 4     ITA   2015    25.14996      NA        NA
## 5     ESP   2015    40.95978      NA        NA
## 6     RUS    NA        NA    2015     1.843
```

- Any variable from A (B) other than those stated in `by=` will be NA for observations that are only in B (A)
- This holds when a variable that is not mentioned in the `by=` argument appears in both datasets:
 - In that case, R adds a data-specific suffix to the names and keeps them both
 - The variable from B (here `year.y`) will be NA for observations that are only in A only (here GBR, ITA, ESP)



2. Merge and reshape

2.1. Merge and append data: example

```
attainment %>% left_join(read.csv("spending.csv"), by = "country")
```

```
##   country year.x share_tertiary year.y share_gdp
## 1    FRA   2015    44.68760    2015    3.398
## 2    GBR   2015    49.94341     NA     NA
## 3    USA   2015    46.51771    2015    3.207
## 4    ITA   2015    25.14996     NA     NA
## 5    ESP   2015    40.95978     NA     NA
```

```
attainment %>% right_join(read.csv("spending.csv"), by = "country")
```

```
##   country year.x share_tertiary year.y share_gdp
## 1    FRA   2015    44.68760    2015    3.398
## 2    USA   2015    46.51771    2015    3.207
## 3    RUS    NA         NA    2015    1.843
```

→ What would be the result of an `inner_join()` here?



2. Merge and reshape

2.2. Reshape data

- It is important to be able to **switch from** the **long** to the **wide** format and conversely
 - Some computations should be done in one format or the other

Wide format

country	year	share_tertiary	share_gdp
FRA	2015	44.69	3.40
USA	2015	46.52	3.21

Long format

country	year	Variable	Value
FRA	2015	share_tertiary	44.69
FRA	2015	share_gdp	3.40
USA	2015	share_tertiary	46.52
USA	2015	share_gdp	3.21



2. Merge and reshape

2.2. Reshape data: From wide to long with pivot_longer()

```
wide <- attainment %>%  
  inner_join(read.csv("spending.csv") %>% select(-year),  
             by = "country")  
wide
```

```
##   country year share_tertiary share_gdp  
## 1     FRA 2015      44.68760      3.398  
## 2     USA 2015      46.51771      3.207
```

→ Pivoting to **long format** can be seen as putting **variables on top of each other** rather side to side

- We need to indicate:
 - **Which variables to stack**
 - The **name of** the variable in which we want the **values** of the stacked variables to be stored
 - The **name of** the variable that will indicate to which **variable** corresponds each value



2. Merge and reshape

2.2. Reshape data: From wide to long with pivot_longer()

```
long <- wide %>%  
  # Which variable to should be stacked  
  pivot_longer(c(share_tertiary, share_gdp),  
    # Where their values should be stored  
    values_to = "Value",  
    # Where to store which variable corresponds each value  
    names_to = "Variable")  
long
```

```
## # A tibble: 4 x 4  
##   country  year Variable      Value  
##   <chr>    <int> <chr>         <dbl>  
## 1 FRA      2015 share_tertiary 44.7  
## 2 FRA      2015 share_gdp      3.40  
## 3 USA      2015 share_tertiary 46.5  
## 4 USA      2015 share_gdp      3.21
```



2. Merge and reshape

2.2. Reshape data: From long to wide with pivot_wider()

- To **pivot in a wide** format we need to indicate:
 - **Which variable** contains **values** of the variables we want to put side to side
 - **Which variable** indicates which **variable** correspond to each value

```
wide <- long %>%  
  # Where the values are  
  pivot_wider(values_from = "Value",  
    # Where the corresponding variable names are  
    names_from = "Variable")  
wide
```

```
## # A tibble: 2 x 4  
##   country  year share_tertiary share_gdp  
##   <chr>    <int>         <dbl>     <dbl>  
## 1 FRA      2015          44.7       3.40  
## 2 USA      2015          46.5       3.21
```


Practice

10:00

- 1) From the `fb` dataset, create a variable `league` equal to `"ligue1"` and a variable `season` equal to `"2021-2022"` and save this new data in an object named `full_fb`
- 2) In `data.zip` you will find the rest of the data for the seasons 2019-2020 to 2021-2022 for the league 1, the bundesliga and the premier league. Append all these data to `full_fb`. Make sure to create the variables `league` and `season` in each data set before appending.
- 3) Use the `separate` function from `tidyr` to extract the number of goals scored by the home and away team
- 4) Convert these variables as numeric and create a variable equal to the sum of the goals from the two teams
- 5) Summarise your data into the total number of goals score per league/season
- 6) Reshape your data such that you have 1 row per league and 1 column per season

You've got 10 minutes!

Solution

1) From the `fb` dataset, create a variable `league` equal to `"ligue1"` and a variable `season` equal to `"2021-2022"` and save this new data in an object named `full_fb`

```
full_fb <- fb %>% mutate(league = "ligue1", season = "2021-2022")
```

2) In `data.zip` you will find the rest of the data for the seasons 2019-2020 to 2021-2022 for the league 1, the bundesliga and the premier league. Append all these data to `full_fb`. Make sure to create the variables `league` and `season` in each data set before appending.

```
full_fb <- full_fb %>%  
  bind_rows(read.csv("ligue1_2021.csv") %>% mutate(league = "ligue1", season = "2020-2021")) %>%  
  bind_rows(read.csv("ligue1_1920.csv") %>% mutate(league = "ligue1", season = "2019-2020")) %>%  
  bind_rows(read.csv("preml_2122.csv") %>% mutate(league = "preml", season = "2021-2022")) %>%  
  bind_rows(read.csv("preml_2021.csv") %>% mutate(league = "preml", season = "2020-2021")) %>%  
  bind_rows(read.csv("preml_1920.csv") %>% mutate(league = "preml", season = "2019-2020")) %>%  
  bind_rows(read.csv("bundes_2122.csv") %>% mutate(league = "bundes", season = "2021-2022")) %>%  
  bind_rows(read.csv("bundes_2021.csv") %>% mutate(league = "bundes", season = "2020-2021")) %>%  
  bind_rows(read.csv("bundes_1920.csv") %>% mutate(league = "bundes", season = "2019-2020"))
```

Solution

3) Use the **separate** function from **tidyr** to extract the number of goals scored by the home and away team

```
full_fb <- full_fb %>%  
  separate(Score, c("home_score", "away_score"), "-")
```

4) Convert these variables as numeric and create a variable equal to the sum of the goals from the two teams

```
full_fb <- full_fb %>%  
  mutate(home_score = as.numeric(home_score),  
         away_score = as.numeric(away_score),  
         goals = home_score + away_score)
```

5) Summarise your data into the total number of goals score per league/season

```
full_fb <- full_fb %>%  
  group_by(league, season) %>%  
  summarise(goals = sum(goals))
```

Solution

6) Reshape your data such that you have 1 row per league and 1 column per season

```
full_fb %>%  
  pivot_wider(names_from = "season", values_from = "goals")
```

```
## # A tibble: 3 x 4  
## # Groups:   league [3]  
##   league `2019-2020` `2020-2021` `2021-2022`  
##   <chr>      <dbl>      <dbl>      <dbl>  
## 1 bundes      982        928        954  
## 2 ligue1       704       1049       1067  
## 3 preml       1034       1024       1071
```



Overview

1. The dplyr package ✓

- 1.1. Packages
- 1.2. Basic functions
- 1.3. `group_by()` and `summarise()`

2. Merge and reshape ✓

- 2.1. Merge and append data
- 2.2. Reshape data

3. A few words on learning R

- 3.1. When it doesn't work the way you want
- 3.2. Where to find help
- 3.3. When it doesn't work at all

4. Wrap up!



Overview

1. The dplyr package ✓

- 1.1. Packages
- 1.2. Basic functions
- 1.3. group_by() and summarise()

2. Merge and reshape ✓

- 2.1. Merge and append data
- 2.2. Reshape data

3. A few words on learning R

- 3.1. When it doesn't work the way you want
- 3.2. Where to find help
- 3.3. When it doesn't work at all



3. A few words on learning R

3.1. When it doesn't work the way you want

- When things do not work the way you want, **NAs are the usual suspects**
 - For instance, this is how the mean function reacts to NAs:

```
mean(c(1, 2, NA))
```

```
## [1] NA
```

```
mean(c(1, 2, NA), na.rm = T)
```

```
## [1] 1.5
```

- You should systematically **check for NAs!**

```
is.na(c(1, 2, NA))
```

```
## [1] FALSE FALSE TRUE
```



3. A few words on learning R

3.1. When it doesn't work the way you want

- **Don't pipe blindly!**
 - **Check** that each command does what it's expected to do
 - View or print your data **at each step**

```
fb %>%  
  select(Home, Score, Away) %>%  
  head(1)
```

```
##      Home Score  Away  
## 1 Monaco   1-1 Nantes
```

```
fb %>%  
  select(Home, Score, Away) %>%  
  separate(Score, c("home_score", "away_score"), "-") %>%  
  head(1)
```

```
##      Home home_score away_score  Away  
## 1 Monaco           1           1 Nantes
```




3. A few words on learning R

3.2. Where to find help

- Oftentimes things don't work either because:
 - **You don't understand** a function's argument
 - Or **you don't know** that there exists an argument that you should use
- This is precisely what **help files** are made for
 - Every function has a help file, just enter **?** and the name of your **function** in the console
 - The help file will **pop up in the Help tab** of R studio

```
?paste
```

Arguments

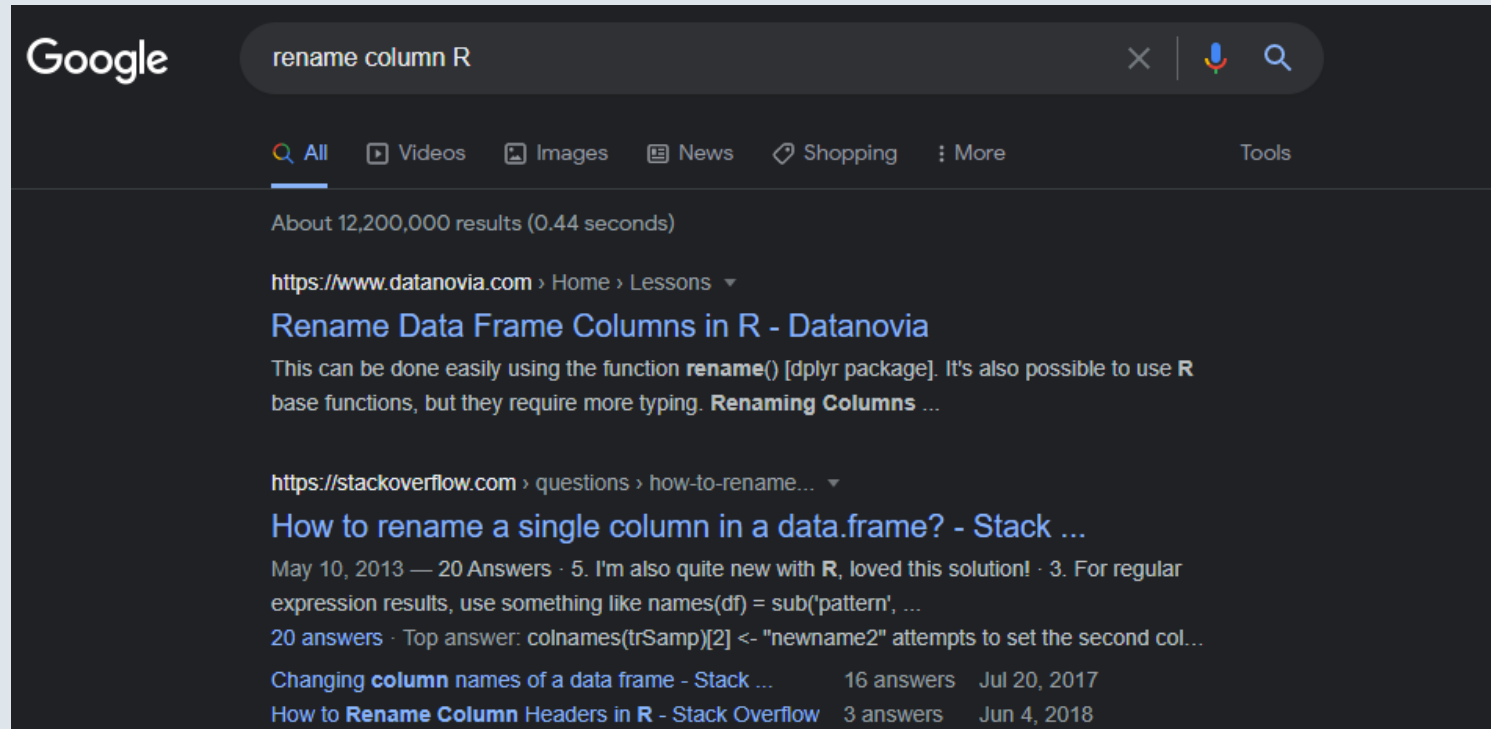
```
...      one or more R objects, to be converted to character vectors.
sep      a character string to separate the terms. Not NA_character_.
collapse an optional character string to separate the results. Not NA_character_.
recycle0 logical indicating if zero-length character arguments should lead to the zero-length character(0)
         after the sep-phase (which turns into "" in the collapse-phase, i.e., when collapse is not NULL).
```



3. A few words on learning R

3.2. Where to find help

- Search on the internet!
 - Your question is for sure already asked and answered on [stackoverflow](#)





3. A few words on learning R

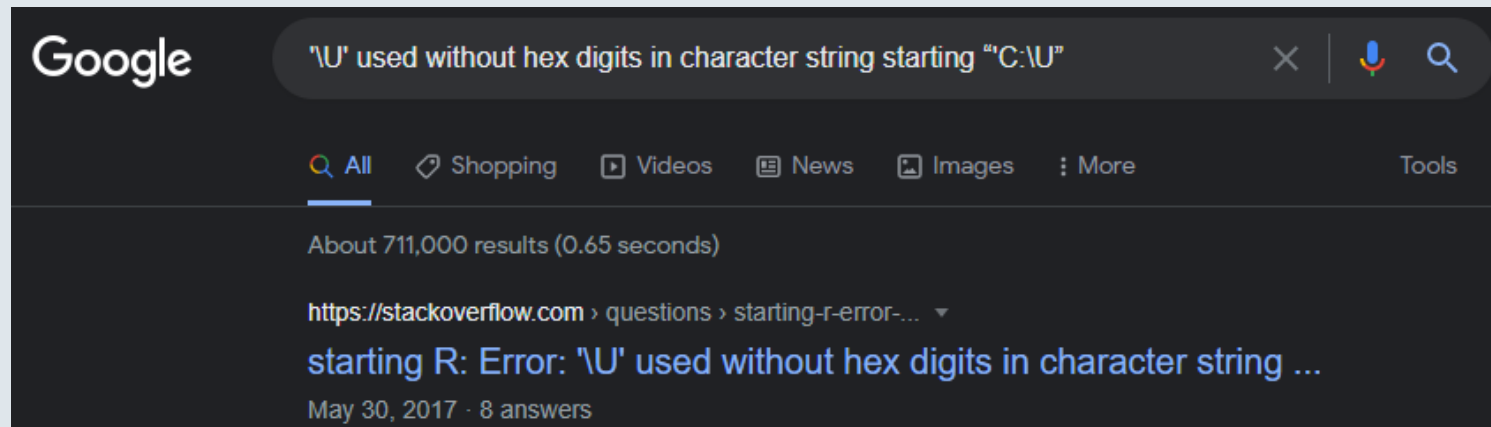
3.3. When it doesn't work at all

- Sometimes R breaks and returns an **error** (usually kind of cryptic)

```
read.csv("C:\\Users\\Documents\\R")
```

```
## Error: '\U' used without hex digits in character string starting ""C:\\U"
```

1. Look for **keywords** that might help you understand where it comes from
2. Paste in on **Google** with the name of your command





Overview

1. The dplyr package ✓

- 1.1. Packages
- 1.2. Basic functions
- 1.3. `group_by()` and `summarise()`

2. Merge and reshape ✓

- 2.1. Merge and append data
- 2.2. Reshape data

3. A few words on learning R ✓

- 3.1. When it doesn't work the way you want
- 3.2. Where to find help
- 3.3. When it doesn't work at all

4. Wrap up!



4. Wrap up!

1. Packages

```
library(dplyr)
```

2. Main dplyr functions

Function	Meaning
mutate()	Modify or create a variable
select()	Keep a subset of variables
filter()	Keep a subset of observations
arrange()	Sort the data
group_by()	Group the data
summarise()	Summarizes variables into 1 observation per group





4. Wrap up!

3. Merge data

```
a <- data.frame(x = c(1, 2, 3), y = c("a", "b", "c"))  
b <- data.frame(x = c(4, 5, 6), y = c("d", "e", "f"))  
c <- data.frame(x = 1:6, z = c("alpha", "bravo", "charlie", "delta", "echo", "foxtrot"))
```

```
a %>% bind_rows(b) %>% left_join(c, by = "x")
```

x	y	z
1	a	alpha
2	b	bravo
3	c	charlie
4	d	delta
5	e	echo
6	f	foxtrot



4. Wrap up!

4. Reshape data

country	year	share_tertiary	share_gdp
FRA	2015	44.69	3.40
USA	2015	46.52	3.21

```
data %>% pivot_longer(c(share_tertiary, share_gdp), names_to = "Variable", values_to = "Value")
```

country	year	Variable	Value
FRA	2015	share_tertiary	44.69
FRA	2015	share_gdp	3.40
USA	2015	share_tertiary	46.52
USA	2015	share_gdp	3.21



For next time

Install the R packages needed for Part I of the course:

ggplot2

rmarkdown

knitr

DT